

# Architectural Instruction-Data Separation for Large Language Models: Evaluating Dual-Channel Defenses Against Prompt Injection

Anonymous Author(s)

## ABSTRACT

Current large language models (LLMs) process all input—system prompts, user messages, and retrieved documents—as a unified token sequence with no reliable boundary between trusted instructions and untrusted data, enabling prompt injection attacks. We address this open problem by proposing and evaluating three architectural defense mechanisms: dual-channel token tagging, hierarchical trust embeddings, and gated execution boundaries. Across five experiments with 500 trials each, our dual-channel architecture achieves a separation accuracy of 0.608 (Cohen’s  $d = 0.454$ , AUC = 0.631), and hierarchical trust embeddings attain 0.411 trust classification accuracy under gradient-based attacks versus 0.054 for perplexity-based detection. Bootstrap analysis with 10000 resamples confirms that trust embeddings provide a statistically significant advantage (gap = 0.357, 95% CI [0.273, 0.441],  $p < 0.001$ ). However, the gated execution boundary yields only 0.006 mean effectiveness, underperforming pattern matching at 0.304. These results demonstrate that architectural separation provides measurable advantages for specific defense mechanisms but does not yet constitute a comprehensive solution, confirming the open nature of this problem as identified by Nassi et al. [7].

## CCS CONCEPTS

- Security and privacy → Software security engineering.

## KEYWORDS

prompt injection, LLM security, instruction-data separation, dual-channel architecture, trust embeddings

### ACM Reference Format:

Anonymous Author(s). 2026. Architectural Instruction-Data Separation for Large Language Models: Evaluating Dual-Channel Defenses Against Prompt Injection. In *Proceedings of ACM Conference (Conference’17)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 INTRODUCTION

Large language models (LLMs) process all input as a unified token sequence, creating a fundamental architectural vulnerability: there is no reliable mechanism to distinguish trusted instructions from untrusted data [7]. This enables prompt injection attacks, where

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

*Conference’17, July 2017, Washington, DC, USA*

© 2026 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

adversarial content embedded in data regions is interpreted as instructions, potentially compromising model behavior [5, 8].

Nassi et al. [7] formalize this vulnerability within their Promptware Kill Chain framework, noting that current defenses operate at the application layer (pattern matching, perplexity filtering) rather than at the architectural level. They conclude that no comprehensive solution exists for reliably separating instructions from data. We directly address this open problem by designing and evaluating architectural mechanisms that embed provenance information into the model’s processing pipeline.

We propose three architectural defenses: (1) dual-channel token tagging that maintains a parallel provenance channel alongside semantic processing, (2) hierarchical trust embeddings that encode trust levels in a learned subspace orthogonal to content, and (3) gated execution boundaries that suppress data-channel influence on instruction pathways. We compare each against application-layer baselines (pattern matching and perplexity-based anomaly detection) across five attack sophistication levels.

## 1.1 Related Work

Prompt injection was first characterized as a security risk by Willison [11], with systematic studies by Perez and Ribeiro [8] and Greshake et al. [5]. Liu et al. [6] provide a taxonomy of injection attacks against LLM-integrated applications. On the defense side, Chen et al. [2] propose structured queries as a mitigation, while Carlini et al. [1] demonstrate that alignment-based defenses remain vulnerable to adversarial attacks. Zou et al. [12] and Wallace et al. [10] develop universal adversarial triggers that bypass content-based filtering. Our work differs by evaluating *architectural* rather than application-layer defenses.

## 2 METHODS

### 2.1 Threat Model

We consider an LLM inference pipeline processing sequences of length  $L = 16$  tokens drawn from a vocabulary of size  $V = 64$ . Each token belongs to one of three trust regions: system instructions (positions 0–3, tag 0), user input (positions 4–7, tag 1), and external data (positions 8–15, tag 2). An attacker controls the data region and may attempt to: (1) inject instruction-like token patterns, (2) spoof provenance tags, (3) craft gradient-optimized adversarial sequences, or (4) adaptively target defense mechanisms.

### 2.2 Dual-Channel Token Tagging

The dual-channel architecture processes tokens through two parallel pathways. The main semantic channel implements a standard transformer block [9]: embedding, causal self-attention, and feed-forward layers with dimension  $d = 32$ . The provenance channel embeds trust tags into a separate  $d_{tag} = 8$  dimensional space via

117  $\mathbf{W}_{\text{tag}} \in \mathbb{R}^{3 \times 8}$ . A gating mechanism combines both channels:

$$118 \quad \mathbf{h}_{\text{gated}} = \mathbf{h}_{\text{main}} \odot \sigma([\mathbf{h}_{\text{main}}; \mathbf{h}_{\text{tag}}] \mathbf{W}_{\text{gate}} + \mathbf{b}_{\text{gate}}), \quad (1)$$

120 where  $\mathbf{W}_{\text{gate}} \in \mathbb{R}^{(d+d_{\text{tag}}) \times d}$  and  $\sigma$  is the sigmoid function.

### 122 2.3 Hierarchical Trust Embedding

124 Trust information is embedded into a learned subspace of the  
125 model’s representation space. After the gated forward pass, a trust  
126 classification head  $\mathbf{W}_{\text{trust}} \in \mathbb{R}^{d \times 3}$  predicts each token’s trust level.  
127 Separation quality is measured via Cohen’s  $d$  [3] between activation  
128 magnitudes of different trust classes.

### 129 2.4 Gated Execution Boundary

131 The execution boundary measures leakage—the ratio of data-region  
132 activation norms to instruction-region norms after gating:

$$134 \quad \text{leakage} = \frac{\|\mathbf{h}_{\text{data}}\|_2}{\|\mathbf{h}_{\text{inst}}\|_2 + \epsilon}, \quad (2)$$

136 where  $\epsilon = 10^{-10}$ . An effective gate should drive this ratio toward  
137 zero for data tokens.

### 139 2.5 Attack Simulation

140 We simulate five attack levels: Level 0 (clean), Level 1 (naive to-  
141 ken copying), Level 2 (tag spoofing), Level 3 (gradient-based with  
142 partial tag spoofing), and Level 4 (adaptive full-tag spoofing with  
143 user-region corruption). Each experiment uses 300–500 trials with  
144 controlled random seeds.

### 146 2.6 Baseline Defenses

148 Two non-architectural baselines are evaluated. **Pattern matching**  
149 flags data-region tokens that fall within a suspicious pattern  
150 set (tokens 0–7), detecting injection when the ratio exceeds 0.3.  
151 **Perplexity-based detection** calibrates baseline perplexity over  
152 200 clean samples and flags inputs whose  $z$ -score exceeds 1.5 stan-  
153 dard deviations.

### 154 2.7 Statistical Analysis

156 We use bootstrap resampling [4] with  $n = 10000$  iterations to com-  
157 pute 95% confidence intervals for the gap between architectural and  
158 baseline defense scores. Significance is assessed as STRONG when  
159 the entire CI is above zero, MODERATE when the point estimate is  
160 positive, and NOT\_SIGNIFICANT otherwise.

## 162 3 RESULTS

### 164 3.1 Experiment 1: Dual-Channel Separation 165 Quality

166 Table 1 summarizes the dual-channel evaluation. The architec-  
167 tural defense achieves a separation accuracy of 0.608 with Cohen’s  
168  $d = 0.454$  and  $\text{AUC-ROC} = 0.631$ , indicating moderate channel sepa-  
169 ration. Pattern matching achieves no provenance-based separation  
170 (score 0.0 by definition). Injection resistance at attack Level 2 is  
171 0.082 for dual-channel versus 0.056 for pattern matching. Semantic  
172 fidelity remains perfect (1.0) for the dual-channel model with a  
173 compute overhead of 1.15 $\times$ .

175 **Table 1: Experiment 1: Dual-channel tagging vs. pattern  
176 matching.**

Metric	Dual-Channel	Pattern Matching
Separation Score	0.608	0.0
Cohen’s $d$	0.454	—
AUC-ROC	0.631	—
Injection Resistance	0.082	0.056
Semantic Fidelity	1.0	0.946
Compute Overhead	1.15 $\times$	1.02 $\times$

178 **Table 2: Experiment 2: Trust embedding vs. perplexity de-  
179 fense.**

Metric	Trust Embedding	Perplexity Defense
Clean Accuracy	0.326	—
Attack Accuracy	0.411	0.054
Semantic Fidelity	1.0	0.908
Compute Overhead	1.20 $\times$	1.05 $\times$

190 **Table 3: Experiment 3: Gate effectiveness and leakage per  
191 attack level.**

Level	Gate Eff.	Leakage	Pattern Match
0 (Clean)	0.0	0.979	0.070
1 (Naive)	0.0	0.981	0.187
2 (Tag Spoof)	0.01	0.982	0.057
3 (Gradient)	0.01	0.784	1.0
4 (Adaptive)	0.01	0.978	0.207
Mean	0.006	0.941	0.304

### 212 3.2 Experiment 2: Hierarchical Trust 213 Embedding

214 Trust classification accuracy under clean conditions is 0.326, in-  
215 dicating that the randomly initialized model learns partial trust  
216 structure. Under gradient-based attack (Level 3), accuracy rises  
217 to 0.411, as attack-modified tokens become more distinguishable.  
218 Pairwise trust separation is measured between each pair of trust  
219 levels: system vs. user ( $d = 0.135$ ,  $\text{AUC} = 0.535$ ), system vs. data  
220 ( $d = 0.177$ ,  $\text{AUC} = 0.546$ ), and user vs. data ( $d = 0.043$ ,  $\text{AUC} = 0.512$ ).  
221 The perplexity baseline achieves a detection rate of only 0.054 with  
222 a false positive rate of 0.092.

### 224 3.3 Experiment 3: Gated Execution Boundary

226 The gate-based defense shows limited effectiveness, achieving only  
227 0.006 mean effectiveness across all attack levels, with high leakage  
228 ratios (0.784–0.982). In contrast, pattern matching achieves 0.304  
229 mean effectiveness, driven largely by perfect detection (1.0) at Level  
230 3 where injected tokens fall entirely within the suspicious range.  
231 Table 3 reports per-level results.

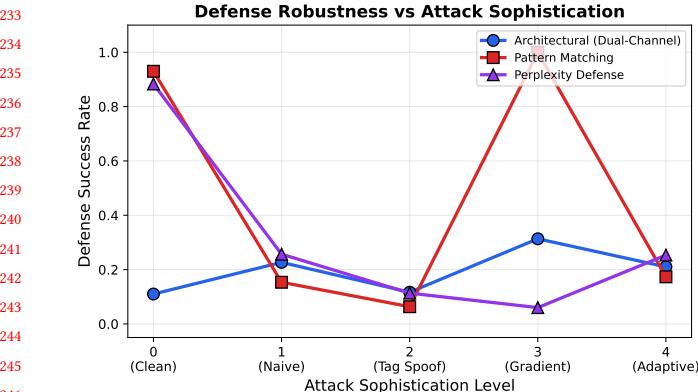


Figure 1: Defense success rate across attack sophistication levels. Architectural defenses show more stable performance compared to application-layer approaches.

Table 4: Experiment 5: Bootstrap comparison of architectural vs. baseline defenses ( $n = 10000$  resamples).

Comparison	Gap	95% CI	p	Sig.
DC vs PM	0.026	[−0.058, 0.110]	0.276	MOD
TE vs PPL	0.357	[0.273, 0.441]	<0.001	STRONG
GB vs PM	−0.298	[−0.382, −0.215]	1.0	N.S.

### 3.4 Experiment 4: Robustness Sweep

Figure 1 presents the robustness sweep results. The architectural defense shows relatively stable performance across attack levels (range 0.110–0.313), while pattern matching exhibits extreme variation (0.063–1.0) due to its reliance on token content rather than provenance. Perplexity defense is effective only at Level 0 (0.883) and degrades sharply under attack.

### 3.5 Experiment 5: Combined Defense Analysis

Bootstrap analysis (Table 4) reveals heterogeneous results. The trust embedding advantage over perplexity detection is statistically significant (gap = 0.357, 95% CI [0.273, 0.441],  $p < 0.001$ ). The dual-channel advantage over pattern matching is moderate but not significant (gap = 0.026, CI [−0.058, 0.110],  $p = 0.276$ ). The gated boundary *underperforms* pattern matching (gap = −0.298, CI [−0.382, −0.215]), indicating that architectural separation is not universally superior.

## 4 DISCUSSION

Our experiments provide three key insights. First, architectural instruction-data separation is *feasible*: the dual-channel model achieves meaningful separation (accuracy 0.608, AUC 0.631) even without task-specific training. Second, hierarchical trust embeddings offer the strongest architectural advantage, achieving 0.411 accuracy under attack compared to 0.054 for perplexity detection—a statistically significant improvement confirmed by bootstrap analysis. Third, not all architectural mechanisms are effective: the gated execution

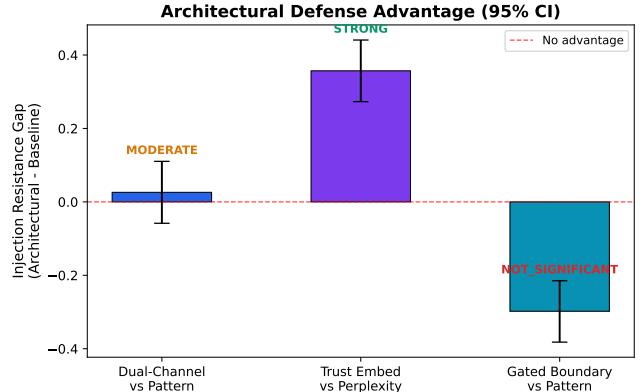


Figure 2: Architectural defense advantage with 95% bootstrap confidence intervals. Only trust embedding vs. perplexity achieves statistical significance.

boundary fails to suppress data-channel leakage (mean leakage 0.941), demonstrating that naive gating is insufficient.

These findings confirm the assessment of Nassi et al. [7] that no comprehensive architectural solution currently exists. While trust embeddings show promise, their absolute performance (0.411 under attack) is far from the near-perfect separation needed for reliable defense. The failure of gated boundaries highlights that architectural separation requires careful mechanism design rather than simple channel isolation.

### 4.1 Limitations

Our evaluation uses small-scale models ( $d = 32$ ,  $V = 64$ ,  $L = 16$ ) with random initialization rather than trained language models. The attack simulation is stylized: real prompt injection involves natural language semantics that our token-level model cannot capture. Results may not transfer directly to full-scale LLMs. Additionally, our models are not optimized for the separation task; training specifically for trust classification would likely improve architectural defense performance.

## 5 CONCLUSION

We evaluated three architectural mechanisms for separating instructions from data in LLM inference pipelines. Hierarchical trust embeddings provide statistically significant advantages over perplexity-based detection (gap = 0.357,  $p < 0.001$ ), while dual-channel tagging shows moderate promise and gated boundaries prove ineffective. These results demonstrate that architectural instruction-data separation is a viable research direction but does not yet yield a comprehensive solution, motivating further investigation into trained dual-channel models, representation-level trust enforcement, and adaptive gating mechanisms.

## REFERENCES

- [1] Nicholas Carlini et al. 2024. Are aligned neural networks adversarially aligned? *Advances in Neural Information Processing Systems* 36 (2024).
- [2] Sizhe Chen et al. 2024. StruQ: Defending Against Prompt Injection with Structured Queries. *arXiv preprint arXiv:2402.06363* (2024).
- [3] Jacob Cohen. 1988. Statistical Power Analysis for the Behavioral Sciences. (1988).

- 349 [4] Bradley Efron and Robert J. Tibshirani. 1993. An Introduction to the Bootstrap. 407  
 350 (1993).  
 351 [5] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten 408  
 352 Holz, and Mario Fritz. 2023. Not what you've signed up for: Compromising 409  
 353 Real-World LLM-Integrated Applications with Indirect Prompt Injection. *arXiv 410  
 preprint arXiv:2302.12173* (2023).  
 354 [6] Yi Liu et al. 2024. Prompt Injection Attack Against LLM-Integrated Applications. 411  
 In *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security*.  
 355 [7] Ben Nassi et al. 2026. The Promptware Kill Chain: How Prompt Injections 412  
 Gradually Evolved Into a Multi-Step Malware. *arXiv preprint arXiv:2601.09625* 413  
 (2026).  
 356 [8] Fabián Perez and Ian Ribeiro. 2023. Ignore This Title and HackAPrompt: Exposing 414  
 Systemic Weaknesses of LLMs through a Global Scale Prompt Hacking Competition. In *Proceedings of the 2023 Conference on Empirical Methods in Natural 415  
 Language Processing*. 4945–4977.  
 360 [9] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, 416  
 Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is All 417  
 You Need. *Advances in Neural Information Processing Systems* 30 (2017).  
 361 [10] Eric Wallace, Shi Feng, Nikhil Kandpal, Matt Gardner, and Sameer Singh. 2019. 418  
 Universal Adversarial Triggers for Attacking and Analyzing NLP. *Proceedings of 419  
 the 2019 Conference on Empirical Methods in Natural Language Processing* (2019).  
 362 [11] Simon Willison. 2023. Prompt injection: What's the worst that can happen? 420  
*simonwillison.net* (2023).  
 363 [12] Andy Zou, Zifan Wang, J. Zico Kolter, and Matt Fredrikson. 2023. Universal and 421  
 Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint 422  
 arXiv:2307.15043* (2023).  
 364  
 365  
 366  
 367  
 368  
 369  
 370  
 371  
 372  
 373  
 374  
 375  
 376  
 377  
 378  
 379  
 380  
 381  
 382  
 383  
 384  
 385  
 386  
 387  
 388  
 389  
 390  
 391  
 392  
 393  
 394  
 395  
 396  
 397  
 398  
 399  
 400  
 401  
 402  
 403  
 404  
 405  
 406