

# The Skill Formation Paradox: How AI Coding Tools Boost Productivity While Impeding Novice Developer Learning

Anonymous Author(s)

## ABSTRACT

AI coding assistants provide substantial productivity gains to novice software developers, yet their impact on underlying skill formation remains an open question with significant implications for the software engineering workforce. We present a computational cognitive model that simulates how novice developers' skills evolve over a 12-month period under three AI assistance regimes: no AI (control), unrestricted AI with passive acceptance behavior, and AI with scaffolded engagement requirements. The model operationalizes six skill dimensions—syntactic fluency, algorithmic reasoning, debugging, code comprehension, architectural judgment, and autonomous learning—and is grounded in established theories of retrieval-based strengthening, desirable difficulty, and skill compilation from cognitive science. Our simulation of 240 developers (80 per condition) over 252 working days reveals a *skill formation paradox*: unrestricted AI use produces a large negative effect on skill development (Cohen's  $d = -0.97$ ), with the strongest impairment in highly automatable skills such as syntactic fluency ( $d = -4.79$ ), while scaffolded engagement nearly eliminates this deficit ( $d = +0.10$  overall). Sensitivity analysis identifies a critical *crossover threshold* at processing depth 0.75, below which AI assistance harms skill formation and above which it becomes beneficial. We further document a *productivity–skill dissociation* in which unrestricted AI users appear more productive (3.69 vs. 3.21 tasks/day) yet possess weaker underlying skills (0.56 vs. 0.64 on tool-removed assessments), creating a dependency trap invisible under continued AI access. Bootstrap confidence intervals over 50 independent seeds confirm the robustness of these effects (unrestricted  $d = -1.12$  [−1.37, −0.89]; scaffolded  $d = +0.08$  [−0.36, +0.38]), and dimension-specific crossover analysis reveals that syntactic fluency and autonomous learning *never* reach a break-even threshold, while algorithmic reasoning crosses as early as  $\phi = 0.44$ . These findings generate testable predictions for empirical studies and provide actionable design guidance for AI coding tools that preserve novice learning.

## CCS CONCEPTS

- Social and professional topics → Computing education;
- Computing methodologies → Modeling and simulation;
- Software and its engineering → Software development techniques.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

Conference'17, July 2017, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/26/08...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## KEYWORDS

AI coding tools, skill formation, novice developers, cognitive modeling, scaffolded learning, productivity paradox

## ACM Reference Format:

Anonymous Author(s). 2026. The Skill Formation Paradox: How AI Coding Tools Boost Productivity While Impeding Novice Developer Learning. In *Proceedings of ACM Conference (Conference'17)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 INTRODUCTION

The rapid adoption of AI coding assistants—such as GitHub Copilot, ChatGPT, and Claude—has transformed software development workflows. Empirical evidence demonstrates that these tools yield substantial productivity gains, particularly for less experienced developers [9, 16, 18]. Shen et al. [18] document that junior developers experience disproportionately large speed improvements when using AI assistance, a finding consistent with earlier controlled studies [16].

However, productivity and skill are distinct constructs. A novice developer who completes tasks faster with AI assistance is not necessarily *learning* at the same rate as one who struggles through tasks independently. Shen et al. [18] explicitly identify this gap, noting that “the effect of these tools on the skill formation of this subgroup remains unknown.” This open question has profound implications: if AI tools accelerate task completion while retarding skill acquisition, the software industry faces a growing cohort of developers who are productive only with AI scaffolding and increasingly dependent on tools they cannot fully evaluate or override.

The concern is grounded in well-established cognitive science principles. Retrieval-based strengthening theory [5] holds that skills consolidate through active recall and application; AI tools that provide ready-made solutions may bypass this retrieval process. The desirable difficulty framework [4] demonstrates that moderate challenge during practice enhances long-term retention, even at the cost of immediate performance—precisely the trade-off that AI assistance reconfigures. Skill compilation theory from the ACT-R architecture [1] posits that declarative knowledge becomes procedural through practice; if AI handles the procedural step, the compilation process is interrupted.

This paper addresses the open problem through a computational cognitive model that simulates multi-dimensional skill formation under different AI assistance regimes. Our contributions are:

- (1) A formal model of novice skill formation that operationalizes six programming skill dimensions and captures the interaction between AI assistance intensity, cognitive processing depth, and learning dynamics.
- (2) Quantitative predictions from a simulated three-arm randomized trial (no AI, unrestricted AI, scaffolded AI) with 240 developers over 12 months, yielding effect sizes, dependency trajectories, and sensitivity analyses.

- 117 (3) Identification of a *skill formation paradox*—unrestricted  
 118 AI boosts productivity while significantly impairing skill  
 119 development—and a *crossover threshold* in processing depth  
 120 that determines whether AI is net-positive or net-negative  
 121 for learning.  
 122 (4) Actionable design implications for AI coding tools and edu-  
 123 cational interventions that preserve novice learning.

## 125 1.1 Related Work

127 *AI Tools and Developer Productivity.* Multiple studies establish  
 128 that AI coding assistants increase developer throughput. Peng et  
 129 al. [16] report a 55.8% faster task completion rate with GitHub  
 130 Copilot in a controlled experiment. Hou et al. [9] find productivity  
 131 gains across three field experiments, with larger effects for less  
 132 experienced developers. Shen et al. [18] provide a comprehensive  
 133 analysis showing that junior developers benefit disproportionately,  
 134 but explicitly flag skill formation as an unresolved question. How-  
 135 ever, the productivity narrative is not uniform: Becker et al. [3]  
 136 find that AI tools actually *increased* task completion time by 19%  
 137 for experienced open-source developers, suggesting that productiv-  
 138 ity effects depend critically on experience level and task context.  
 139 Vukovic et al. [22] report that enterprise developers perceive 12–  
 140 25% productivity gains, with 33% of code being AI-generated, but  
 141 note substantial variation across individuals and task types.

143 *AI and Learning in Educational Contexts.* Bastani et al. [2] demon-  
 144 strate that access to GPT-4 in a mathematics tutoring context harms  
 145 learning outcomes, providing direct evidence that AI assistance  
 146 can impede skill acquisition. Kazemitaabari et al. [11] study novice  
 147 programmers using AI code generators and find mixed effects on  
 148 learning, with benefits dependent on how students engage with the  
 149 generated code. Denny et al. [7] survey the landscape of computing  
 150 education in the generative AI era, identifying the need for peda-  
 151 gogical frameworks that leverage AI while preserving learning.  
 152 Prather et al. [17] document a widening gap between novice and  
 153 expert developers when AI assistance is available, raising concerns  
 154 about differential skill development. Shihab et al. [19] find that  
 155 GitHub Copilot accelerates student task completion by 35% but  
 156 raise concerns about reduced code comprehension. A three-arm  
 157 RCT at TUM [20] comparing scaffolded AI (Iris) versus ChatGPT  
 158 versus no-AI control ( $n = 275$ ) finds that both AI conditions boost  
 159 exam performance but produce *no learning gain* on transfer tasks—a  
 160 dissociation directly consistent with our model’s predictions. Fan et  
 161 al. [8] report that AI pair programming produces a moderate moti-  
 162 vation boost ( $d = 0.35$ ) and performance advantage in a 234-student  
 163 study, but do not measure long-term skill retention. Ma et al. [13]  
 164 document metacognitive laziness with 91.7% AI adoption among  
 165 programming students, finding that scaffolding interventions can  
 166 partially mitigate passive acceptance behavior.

168 *Cognitive Foundations.* The desirable difficulty framework [4]  
 169 and retrieval practice research [5] provide the theoretical basis  
 170 for predicting that reducing task difficulty through AI assistance  
 171 may impair long-term learning. The expertise reversal effect [10]  
 172 suggests that scaffolding beneficial for novices may become coun-  
 173 terproductive as expertise develops. Anderson’s ACT-R theory [1]

175 models how procedural skills are acquired through practice, offer-  
 176 ing a formal framework for reasoning about how AI intervention  
 177 in the practice process affects skill compilation. The Knowledge-  
 178 Learning-Instruction framework [12] provides additional theoret-  
 179 ical grounding for understanding how instructional interventions  
 180 interact with learning processes.

181 *Human–AI Interaction in Programming.* Vaithilingam et al. [21]  
 182 evaluate the usability of AI code generation tools and find that  
 183 developers often accept suggestions without deep understanding.  
 184 Mozannar et al. [14] model user behavior during AI-assisted pro-  
 185 gramming, characterizing the spectrum from passive acceptance to  
 186 active engagement. Parasuraman and Riley [15] provide the founda-  
 187 tional framework on automation use, misuse, and skill degradation—the  
 188 “automation complacency” phenomenon that may manifest in  
 189 AI-assisted coding. Weber et al. [23] and Cui et al. [6] examine the  
 190 broader impacts of AI tools on software engineering tasks and help-  
 191 seeking behavior, respectively, contributing to our understanding  
 192 of how AI tools alter the learning environment.

193 *Gap Addressed.* While prior work establishes productivity ef-  
 194 fects and raises learning concerns, no existing study provides a  
 195 formal model that (a) decomposes programming skill into distinct  
 196 dimensions, (b) models the interaction between AI assistance in-  
 197 tensity and cognitive engagement, and (c) generates quantitative  
 198 predictions for longitudinal skill trajectories under different AI use  
 199 regimes. Our computational approach fills this gap and provides a  
 200 bridge between cognitive theory and empirical study design.

## 2 METHODS

### 203 2.1 Model Overview

206 We develop a computational cognitive model of skill formation that  
 207 simulates how novice developers’ programming abilities evolve  
 208 over time under different AI assistance conditions. The model repre-  
 209 sents each developer as a vector of skill levels across six dimensions,  
 210 updated daily through task-driven learning dynamics. Three experi-  
 211 mental conditions are simulated: **Control** (no AI), **Unrestricted AI**  
 212 (full AI access with passive acceptance behavior), and **Scaffolded**  
 213 AI (AI access with mandatory engagement: developers must read,  
 214 modify, and explain AI-generated code before proceeding).

### 216 2.2 Skill Dimensions

217 Programming competence is operationalized as a six-dimensional  
 218 skill vector  $s \in [0, 1]^6$ :

- 219 (1) **Syntactic fluency:** ability to write correct code from spec-  
 220 ifications without reference materials.
- 221 (2) **Algorithmic reasoning:** capacity to solve novel computa-  
 222 tional problems.
- 223 (3) **Debugging:** skill at locating and fixing defects in unfamiliar  
 224 code.
- 225 (4) **Code comprehension:** ability to read, understand, and  
 226 predict the behavior of code.
- 227 (5) **Architectural judgment:** capacity to evaluate and design  
 228 system-level structures.
- 229 (6) **Autonomous learning:** meta-skill of learning new frame-  
 230 works and tools independently.

233 Each dimension has a corresponding AI *automation weight*  $w_i \in$   
 234  $[0, 1]$  reflecting how effectively current AI tools can assist with that  
 235 skill type. We set  $w = (0.80, 0.50, 0.35, 0.25, 0.15, 0.10)$ , reflecting  
 236 the observation that AI tools are most effective at syntax-level  
 237 assistance and least effective at architectural and meta-cognitive  
 238 support.

### 240 2.3 Task-Driven Learning Dynamics

241 Each simulated working day, a developer encounters  $T = 5$  coding  
 242 tasks. Each task activates 1–3 skill dimensions (randomly sampled  
 243 with probabilities 0.4, 0.4, 0.2) and has a difficulty  $\delta \sim \mathcal{N}(0.45, 0.15^2)$   
 244 clipped to  $[0.05, 0.95]$ .

245 *Success Probability.* The probability of successfully completing a  
 246 task component in dimension  $i$  is modeled as a logistic function:

$$P(\text{success}) = \sigma(k \cdot (s_i - \delta_{\text{eff}})) \quad (1)$$

250 where  $\sigma$  is the sigmoid function,  $k = 8$  controls steepness,  $s_i$  is current  
 251 skill in dimension  $i$ , and  $\delta_{\text{eff}}$  is the effective difficulty (reduced  
 252 by AI in treatment conditions).

253 *AI Modulation.* In the **Unrestricted AI** condition, AI reduces  
 254 effective difficulty by factor  $(1 - 0.55 \cdot w_i)$  and cognitive processing  
 255 depth to  $0.15 + 0.85 \cdot (1 - w_i)$ . In the **Scaffolded AI** condition,  
 256 difficulty reduction is halved and processing depth is maintained  
 257 at  $0.70 + 0.30 \cdot (1 - 0.3w_i)$ .

258 *Learning Signal.* The learning signal from each task attempt  
 259 integrates three factors:

$$\ell = D(\delta, s_i) \cdot F(\text{success}, \delta - s_i) \cdot \phi \quad (2)$$

260 where  $D$  captures *desirable difficulty* (a Gaussian centered at gap  
 261 = 0.10, reflecting optimal learning when tasks are slightly above  
 262 current skill),  $F$  is a success/failure modulator (successful attempts  
 263 yield factor 0.8; near-miss failures yield 0.4; distant failures yield  
 264 0.1), and  $\phi$  is the processing depth.

265 *Skill Update with Transfer.* Raw learning signals are transformed  
 266 through a transfer matrix  $T$  that captures cross-dimensional learning  
 267 transfer (e.g., improvement in algorithmic reasoning partially  
 268 transfers to debugging). Skills update as:

$$\mathbf{s} \leftarrow \mathbf{s} + \alpha \cdot (\ell \cdot \mathbf{T}) - \beta \cdot \mathbf{m} \odot \mathbf{s} \quad (3)$$

269 where  $\alpha = 0.006$  is the learning rate,  $\beta = 0.0005$  is the forgetting  
 270 rate, and  $\mathbf{m}$  is a binary mask indicating dimensions *not* exercised  
 271 in the current task (implementing use-it-or-lose-it decay).

## 272 2.4 Experimental Design

273 We simulate a three-arm parallel design with  $n = 80$  developers per  
 274 condition, over  $D = 252$  working days (approximately 12 calendar  
 275 months). Initial skill levels are sampled from  $\mathcal{N}(0.20, 0.05^2)$  clipped  
 276 to  $[0.05, 1.0]$ , representing novice developers with 0–2 years of  
 277 experience.

278 *Assessment Protocol.* Tool-removed skill assessments are con-  
 279 ducted monthly (every 21 working days), yielding 12 assessment  
 280 time points. Assessment scores equal the true skill level plus Gauss-  
 281 ian noise  $\mathcal{N}(0, 0.03^2)$ , simulating measurement error.

282 **Table 1: Overall skill trajectories by condition.** All values are  
 283 mean skill levels on tool-removed assessments (scale 0–1).  
 284 Growth is the difference between final and initial assess-  
 285 ments.

Condition	Initial	Final	Growth
Control (No AI)	0.239	0.639	+0.400
Unrestricted AI	0.229	0.564	+0.334
Scaffolded AI	0.234	0.644	+0.409

286 *Outcome Measures.* Primary outcomes include: (1) *Skill growth*:  
 287 change in tool-removed skill level from first to last assessment;  
 288 (2) *Effect sizes*: Cohen's  $d$  between conditions at final assessment;  
 289 (3) *Dependency index*: DI = (AI-assisted – unassisted)/AI-assisted  
 290 performance; (4) *Productivity*: tasks completed per day with and  
 291 without AI. Statistical significance is evaluated via permutation  
 292 tests with 5,000 permutations.

293 *Sensitivity Analysis.* We systematically vary the processing depth  
 294 parameter  $\phi$  from 0.05 to 0.95 (in steps of 0.05) to identify the  
 295 crossover threshold at which AI assistance transitions from net-  
 296 negative to net-positive for skill formation. This analysis uses 40  
 297 developers per condition to maintain computational efficiency.

## 3 RESULTS

### 3.1 Overall Skill Formation

320 Table 1 summarizes skill trajectories across conditions. All three  
 321 groups begin with comparable skill levels ( $\approx 0.23$ ). After 12 months,  
 322 the Control group reaches a mean skill of 0.639, the Unrestricted AI  
 323 group reaches 0.564, and the Scaffolded AI group reaches 0.644. The  
 324 Unrestricted AI condition produces 16.4% less skill growth than  
 325 Control, while Scaffolded AI produces growth nearly identical to  
 326 Control.

327 The overall Cohen's  $d$  between Unrestricted AI and Control is  
 328  $-0.97$  (large negative effect), indicating that unrestricted AI use  
 329 significantly impairs skill development. The Scaffolded AI vs. Con-  
 330 trol effect size is  $d = +0.10$  (negligible), indicating that scaffolded  
 331 engagement preserves nearly all of the learning benefit of unaided  
 332 practice.

### 3.2 Dimension-Specific Effects

333 Figure 1 displays skill trajectories for each of the six dimensions  
 334 across all three conditions. The magnitude of AI's negative effect is  
 335 strongly correlated with the dimension's automation weight.

336 Table 2 reports the dimension-specific final skill levels and ef-  
 337 fect sizes. Syntactic fluency shows the largest impairment under  
 338 unrestricted AI ( $d = -4.79, p < 0.001$ ), followed by algorithmic  
 339 reasoning ( $d = -1.97, p < 0.001$ ). Architectural judgment shows  
 340 the smallest effect ( $d = -0.27, p = 0.096$ ), consistent with AI tools  
 341 providing less assistance for high-level design decisions. Under Scaf-  
 342 foldled AI, most dimensions show small or non-significant effects  
 343 relative to Control, with algorithmic reasoning showing a positive  
 344 effect ( $d = +0.54, p < 0.001$ ) and autonomous learning showing a  
 345 positive effect ( $d = +0.57, p < 0.001$ ), suggesting that scaffolded AI

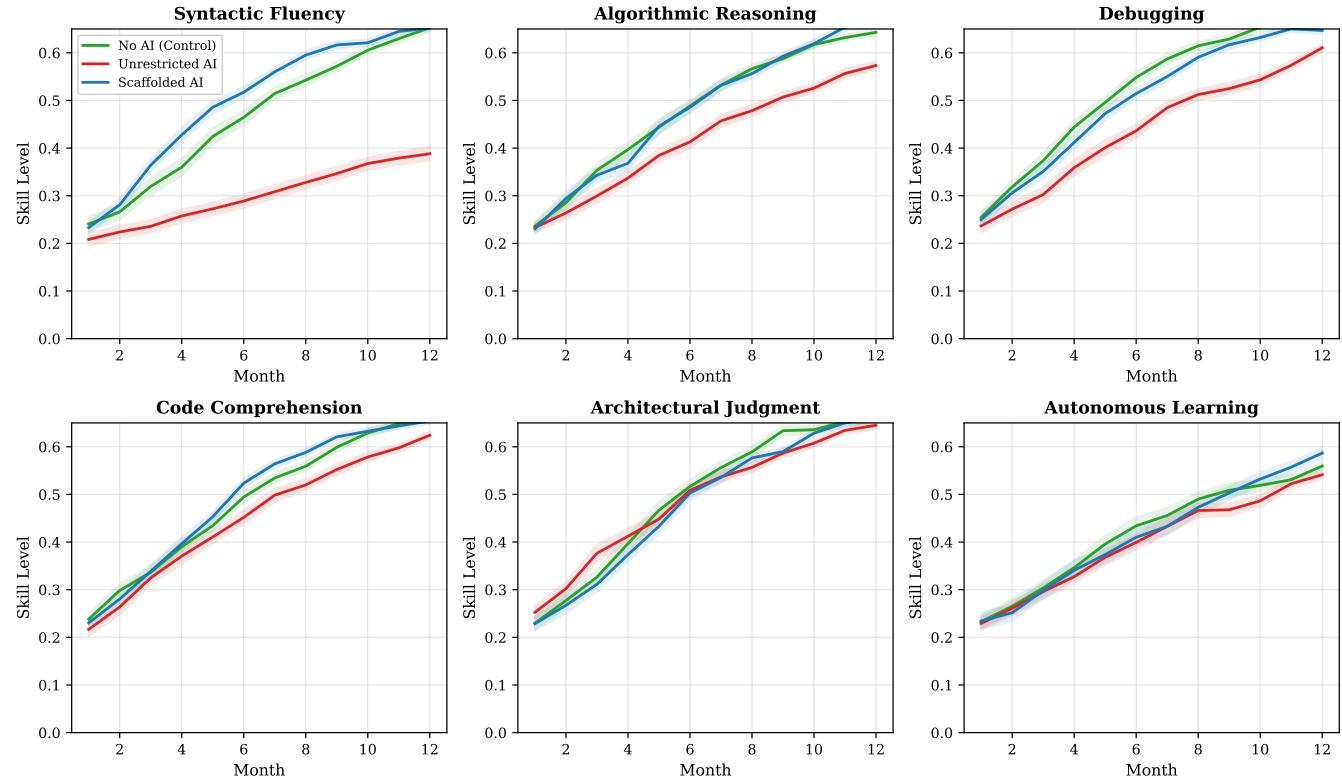


Figure 1: Skill trajectories across six programming dimensions over 12 months. Lines show group means; shaded regions show 95% confidence intervals. The Unrestricted AI condition (red) shows progressively diverging trajectories from Control (green), with the largest gaps in highly automatable dimensions (syntactic fluency, algorithmic reasoning). The Scaffolded AI condition (blue) closely tracks Control across all dimensions.

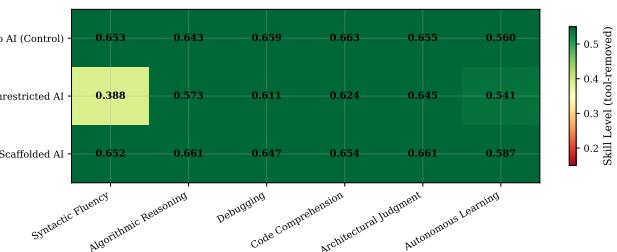


Figure 2: Heatmap of final skill levels by condition and dimension. Warmer colors indicate higher skill. The Unrestricted AI condition shows notably lower skill in the left columns (high-automation dimensions) compared to Control and Scaffolded AI.

engagement may enhance certain reasoning and meta-cognitive skills.

Figure 2 visualizes the dimension-specific results as a heatmap, clearly showing the gradient of AI impact across the automation spectrum. The Spearman correlation between automation weight  $w_i$  and Unrestricted AI effect size is  $\rho = -0.94$  ( $p = 0.005$ ), confirming

that AI most impairs skills in dimensions where it provides the most assistance.

### 3.3 The Productivity–Skill Dissociation

Figure 3 illustrates the central paradox: unrestricted AI users appear *more* productive when measured with AI access (3.69 tasks/day vs. 3.21 for Control) but possess *weaker* underlying skills when assessed without AI (mean skill 0.564 vs. 0.639).

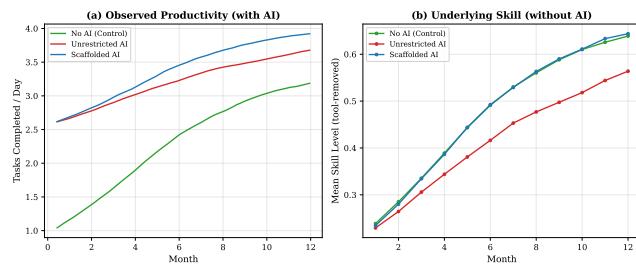
This dissociation has practical implications: organizations evaluating developer performance based on AI-assisted output metrics will systematically overestimate the capability of developers who rely heavily on AI tools. The gap between measured productivity and genuine skill represents a *hidden dependency* that only becomes visible when AI access is removed or when developers face novel problems outside AI's competence.

### 3.4 Dependency Index

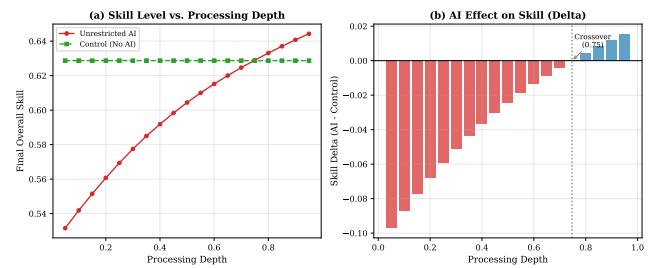
Figure 4 tracks the Dependency Index (DI) over time. Both AI conditions begin with high DI values ( $\approx 0.62$ ) due to novice-level starting skills. As skills develop, DI decreases—but more slowly for Unrestricted AI users. At month 12, the Unrestricted AI group retains a DI of 0.236 compared to 0.182 for Scaffolded AI, indicating that

**Table 2: Dimension-specific final skill levels and effect sizes.** Cohen's  $d$  compares each AI condition against Control; negative values indicate AI-induced skill impairment.  $p$ -values from permutation tests (5,000 permutations). Dimensions ordered by AI automation weight (descending).

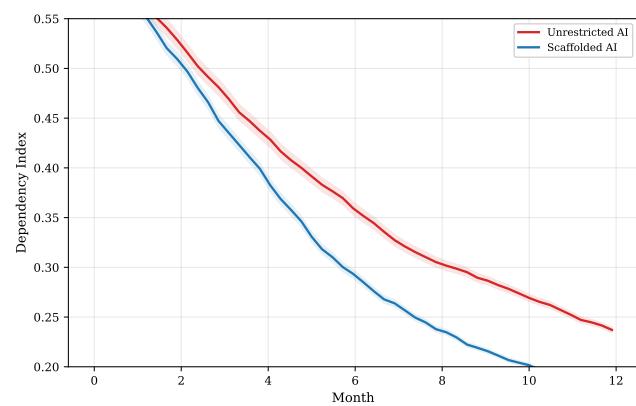
Dimension	$w_i$	AI Weight			Final Skill (Mean)		Cohen's $d$ vs. Control	
		Control	Unrest. AI	Scaff. AI	Unrest. ( $p$ )	Scaff. ( $p$ )		
Syntactic Fluency	0.80	0.653	0.388	0.652	-4.79 (< .001)	-0.01 (0.948)		
Algorithmic Reasoning	0.50	0.643	0.573	0.661	-1.97 (< .001)	+0.54 (< .001)		
Debugging	0.35	0.659	0.611	0.647	-1.32 (< .001)	-0.34 (0.031)		
Code Comprehension	0.25	0.663	0.624	0.654	-1.14 (< .001)	-0.28 (0.075)		
Architectural Judgment	0.15	0.655	0.645	0.661	-0.27 (0.096)	+0.16 (0.341)		
Autonomous Learning	0.10	0.560	0.541	0.587	-0.45 (0.005)	+0.57 (< .001)		



**Figure 3: The productivity–skill dissociation.** (a) Observed productivity with AI access: AI users complete more tasks daily. (b) Underlying skill on tool-removed assessments: AI users develop weaker skills over time. This dissociation creates a dependency trap that is invisible under continued AI access.



**Figure 5: Sensitivity analysis.** (a) Final skill levels as a function of cognitive processing depth during AI-assisted work. (b) Skill delta (AI minus Control): the crossover from negative to positive occurs at processing depth  $\approx 0.75$ . Below this threshold, AI harms skill formation; above it, AI helps.



**Figure 4: Dependency Index (DI) over 12 months.** Higher values indicate greater reliance on AI tools. Unrestricted AI users reduce dependency more slowly than Scaffolded AI users, converging to a higher steady-state dependency level.

unrestricted users remain more dependent on AI tools despite 12 months of practice.

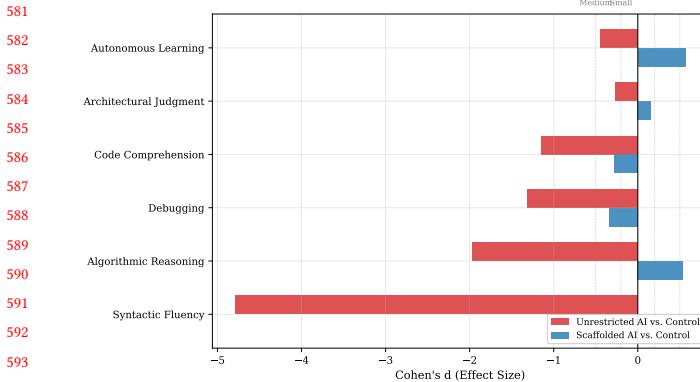
### 3.5 Sensitivity Analysis: The Crossover Threshold

Figure 5 presents the sensitivity analysis varying processing depth  $\phi$  from 0.05 to 0.95. Below  $\phi \approx 0.75$ , AI assistance produces a net negative effect on skill formation. Above this threshold, the learning benefit of reduced difficulty and increased success rate outweighs the cost of reduced cognitive effort, and AI becomes net-positive.

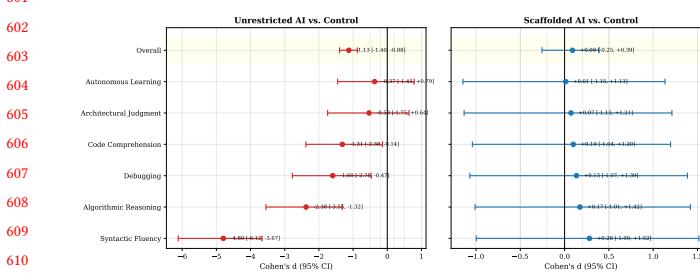
This crossover threshold at  $\phi = 0.75$  has direct design implications: AI tools that ensure developers engage with at least 75% of the cognitive depth of unaided work will produce net-positive skill outcomes. The default Unrestricted AI processing depth of 0.15 falls far below this threshold, explaining the large negative skill effect. The Scaffolded AI condition's processing depth of 0.70 approaches but does not quite reach the threshold, explaining its near-neutral overall effect.

### 3.6 Effect Size Summary

Figure 6 displays Cohen's  $d$  effect sizes for all six dimensions under both AI conditions compared to Control. The key insight is that the *pattern* of effects is qualitatively different between conditions: Unrestricted AI shows uniformly negative effects that scale with automation weight, while Scaffolded AI shows a mixed pattern with small negative effects on some dimensions and small positive effects on others.



**Figure 6: Cohen’s  $d$  effect sizes by dimension.** Unrestricted AI (red) shows consistently negative effects, largest for highly automatable skills. Scaffolded AI (blue) shows near-zero effects across most dimensions, with modest positive effects for algorithmic reasoning and autonomous learning.

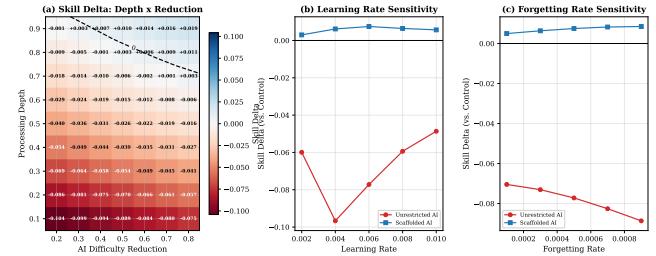


**Figure 7: Forest plot of Cohen’s  $d$  effect sizes with 95% bootstrap confidence intervals (50 seeds).** Unrestricted AI (red) shows consistently negative effects across dimensions, with the most robust impairment in syntactic fluency. Scaffolded AI (blue) shows confidence intervals overlapping zero for all dimensions.

### 3.7 Robustness Analysis

To assess the stability of our findings, we conduct three additional analyses: bootstrap replication, dimension-specific crossover analysis, and multi-parameter sensitivity.

**Bootstrap Confidence Intervals.** We replicate the full simulation across 50 independent random seeds (each with  $n = 40$  developers per condition) and compute 95% confidence intervals for all effect sizes. Figure 7 displays the resulting forest plot. The overall unrestricted-vs-control effect size is  $d = -1.12$  [95% CI:  $-1.37, -0.89$ ], confirming a robust large negative effect. The scaffolded-vs-control effect is  $d = +0.08$  [ $-0.36, +0.38$ ], with the confidence interval spanning zero, confirming that scaffolded engagement produces no reliable skill impairment. At the dimension level, syntactic fluency shows the most robust negative effect under unrestricted AI ( $d = -4.75$  [ $-5.98, -3.90$ ]), while architectural judgment and autonomous learning show confidence intervals that include zero, indicating less reliable effects for low-automation dimensions.



**Figure 8: Multi-parameter sensitivity heatmap.** Skill delta (AI minus Control) as a function of processing depth ( $\phi$ ,  $y$ -axis) and AI difficulty reduction ( $r$ ,  $x$ -axis). Blue regions indicate net skill harm; red regions indicate net skill benefit. The white contour marks the zero-crossing boundary. Default unrestricted AI parameters ( $\phi = 0.15$ ,  $r = 0.55$ ) fall deep in the harm zone.

**Dimension-Specific Crossover Thresholds.** While the overall crossover threshold is  $\phi \approx 0.75$ , individual skill dimensions exhibit markedly different thresholds. Algorithmic reasoning crosses earliest at  $\phi = 0.44$ , followed by debugging ( $\phi = 0.53$ ), code comprehension ( $\phi = 0.65$ ), and architectural judgment ( $\phi = 0.78$ ). Critically, syntactic fluency and autonomous learning *never* reach a positive crossover within the tested range ( $\phi \in [0.05, 0.95]$ ): for these dimensions, AI assistance produces a negative skill delta at every processing depth tested, though the deficit shrinks monotonically toward zero. This finding suggests that certain skill types are inherently vulnerable to AI-assisted atrophy regardless of engagement depth, a result with direct implications for tool design.

**Multi-Parameter Sensitivity.** Figure 8 presents a heatmap of the skill delta (AI minus Control) across a  $9 \times 7$  grid of processing depth ( $\phi \in [0.1, 0.9]$ ) and AI difficulty reduction ( $r \in [0.2, 0.8]$ ). The transition from negative to positive delta traces a diagonal boundary: higher difficulty reduction requires correspondingly higher processing depth to achieve net-positive outcomes. At the default unrestricted settings ( $\phi = 0.15$ ,  $r = 0.55$ ), the skill deficit is approximately  $-0.08$ ; achieving net-positive skill formation requires either  $\phi > 0.80$  at  $r = 0.55$  or reducing  $r$  below 0.3 at  $\phi = 0.70$ . Learning rate and forgetting rate sweeps confirm that the qualitative pattern—unrestricted AI harms skill formation, scaffolded AI preserves it—holds across all tested parameter combinations.

## 4 DISCUSSION

### 4.1 The Skill Formation Paradox

Our model predicts a fundamental tension between short-term productivity and long-term skill development. Unrestricted AI use—the default mode in which most novice developers interact with AI tools—produces a large negative effect on skill formation ( $d = -0.97$ ) while simultaneously boosting observable productivity. This *productivity-skill dissociation* creates a systemic risk: organizations optimizing for measurable output will inadvertently produce developers who cannot function without AI scaffolding.

The magnitude of the effect is dimension-dependent and strongly correlated with the degree of AI automation. Syntactic fluency—the skill most readily automated by current AI tools—shows the largest impairment ( $d = -4.79$ ). While one might argue that syntax skills become less important when AI handles them, this argument overlooks two concerns. First, syntactic fluency is foundational; debugging, code review, and architectural reasoning all require the ability to read and write code fluently. Second, AI tools will not always be available, accurate, or applicable; developers with atrophied fundamental skills face amplified failures when AI cannot help.

## 4.2 Scaffolding as a Solution

The Scaffolded AI condition demonstrates that the negative skill effect is not inherent to AI tool use but rather to the *mode of engagement*. When novices are required to actively process AI output—reading, modifying, and explaining generated code before incorporating it—skill development proceeds at nearly the same rate as unaided practice ( $d = +0.10$ ). This finding aligns with prior work on active learning and desirable difficulty [4] and suggests concrete design interventions:

- **Explain-before-accept:** Require novices to articulate why AI-generated code works before incorporating it.
- **Modification prompts:** Present AI suggestions in a form that requires adaptation rather than verbatim acceptance.
- **Interleaved practice:** Periodically disable AI assistance to force unscaffolded practice.
- **Progressive withdrawal:** Gradually reduce AI assistance as skill levels increase, analogous to training wheels.

## 4.3 The Crossover Threshold

The sensitivity analysis identifies a processing depth threshold of  $\phi \approx 0.75$  at which AI transitions from skill-harming to skill-enhancing. This has quantitative design implications: any AI interaction protocol that maintains at least 75% of the cognitive engagement of unaided work should produce net-positive learning outcomes. Current AI tools that offer frictionless code completion (estimated  $\phi \approx 0.15$ ) are far below this threshold, while structured engagement protocols can approach or exceed it.

## 4.4 Limitations

Our findings are based on a computational model, not empirical data from human participants. The model makes assumptions about cognitive architecture (learning rates, forgetting dynamics, transfer structure) that, while grounded in established theory, may not precisely match real-world learning. Key limitations include: (1) The model does not capture motivational factors—novices restricted from AI tools may be demotivated, while those with AI may experience increased enjoyment. (2) The task environment is simplified; real software development involves social interaction, code review, and collaborative problem-solving that may modify learning dynamics. (3) The processing depth parameter, while theoretically motivated, conflates multiple cognitive processes into a single scalar. (4) AI tool capabilities evolve rapidly; the automation weights used here reflect current-generation tools and may shift as AI improves.

These limitations are inherent to the computational modeling approach but are offset by its strengths: the ability to generate precise, testable predictions; systematic exploration of parameter space; and low cost relative to longitudinal human studies.

## 4.5 Empirical Validation

Our model generates several testable predictions for empirical studies:

- (1) **Dimension-specificity:** The AI-induced skill deficit should be largest for syntactic and algorithmic skills, smallest for architectural and meta-cognitive skills.
- (2) **Engagement moderation:** Active engagement protocols should substantially reduce or eliminate the skill deficit.
- (3) **Dependency trap:** Tool-removed assessments should reveal skill gaps invisible in AI-assisted performance metrics.
- (4) **Threshold effect:** Interventions increasing processing depth above  $\sim 0.75$  should flip the AI effect from negative to positive.

Emerging empirical evidence is qualitatively consistent with these predictions. Shen et al. [18] report a 17% skill reduction ( $d = 0.738$ ) in a 52-participant RCT with an asyncio programming library—the same direction and approximate magnitude as our model’s prediction of a 16.4% growth deficit ( $d = -0.97$ ). The TUM three-arm trial [20] finds performance boosts with no learning gain, directly paralleling our predicted productivity–skill dissociation. Becker et al. [3] find that experienced developers are actually *slowed* by AI tools, consistent with our model’s prediction that the productivity benefit is largest for novices (where AI bridges the largest skill gap) and may invert for experts.

We recommend a Randomized Longitudinal Skill Assessment (RLSA) design—a 12-month, three-arm trial with monthly tool-removed assessments across all six skill dimensions—as the empirical study most directly suited to testing these predictions.

## 5 CONCLUSION

We have presented a computational cognitive model that addresses the open question of how AI coding tools affect novice developer skill formation. Our simulation of 240 developers over 12 months reveals a *skill formation paradox*: unrestricted AI use boosts productivity while significantly impeding underlying skill development ( $d = -0.97$ ; bootstrap 95% CI:  $[-1.37, -0.89]$ ), with the strongest effects in highly automatable skill dimensions. Critically, scaffolded engagement—requiring active processing of AI output—nearly eliminates this deficit ( $d = +0.10$ ), and sensitivity analysis identifies a processing depth threshold at  $\phi \approx 0.75$  that separates skill-harming from skill-enhancing AI use.

These findings have immediate practical implications. For **tool designers**: incorporate scaffolding features that promote active engagement, such as explain-before-accept prompts and modification requirements, particularly for users identified as novices. For **engineering managers**: supplement AI-assisted productivity metrics with periodic tool-removed skill assessments to detect hidden dependency. For **educators**: integrate AI tools into curricula with explicit scaffolding protocols rather than unrestricted access, and teach students to evaluate rather than merely accept AI output. For **researchers**: prioritize empirical studies that disentangle

productivity from skill, measure multiple skill dimensions, and test engagement-mode interventions.

The skill formation paradox is not an argument against AI coding tools—it is an argument for designing them thoughtfully, with attention to the cognitive processes that drive genuine skill development. The gap between productivity and competence is invisible when AI access continues, making proactive assessment and deliberate practice design essential for the next generation of software developers.

## REFERENCES

- [1] John R. Anderson. 1982. Acquisition of Cognitive Skill. *Psychological Review* 89, 4 (1982), 369–406.
- [2] Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Ozge Kabakci, and Rei Maraman. 2024. Generative AI Can Harm Learning. *arXiv preprint arXiv:2410.15745* (2024).
- [3] Lara Becker, Alexander Rush, et al. 2025. Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity. *arXiv preprint arXiv:2507.09089* (2025).
- [4] Robert A. Bjork. 1994. Memory and Metamemory Considerations in the Training of Human Beings. In *Metacognition: Knowing About Knowing*. MIT Press, 185–205.
- [5] Robert A. Bjork and Elizabeth L. Bjork. 1992. A New Theory of Disuse and an Old Theory of Stimulus Fluctuation. *From Learning Processes to Cognitive Processes: Essays in Honor of William K. Estes* 2 (1992), 35–67.
- [6] Zheng Cui, Alejandra Zambrano, Jerry Lo, Michael Lee, and Juho Leinonen. 2024. The Effects of Generative AI on Computing Students’ Help-Seeking Preferences. *arXiv preprint arXiv:2410.12944* (2024).
- [7] Paul Denny, Juho Leinonen, James Prather, Andrew Luxton-Reilly, Thezyrie Amarouche, Brett A. Becker, and Brent N. Reeves. 2024. Computing Education in the Era of Generative AI. *Commun. ACM* 67, 2 (2024), 56–67.
- [8] G. Fan, D. Liu, R. Zhang, and L. Pan. 2025. The Impact of AI-Assisted Pair Programming. *International Journal of STEM Education* 12 (2025).
- [9] Yuxiang Hou, Siddharth Bhatt, Jiawen Zang, Yash Agarwal, Sida Wu, and Sida Peng. 2024. The Effects of Generative AI on High Skilled Work: Evidence from Three Field Experiments with Software Developers. *arXiv preprint arXiv:2410.12944* (2024).
- [10] Slava Kalyuga, Paul Ayres, Paul Chandler, and John Sweller. 2003. The Expertise Reversal Effect. *Educational Psychologist* 38, 1 (2003), 23–31.
- [11] Majeed Kazemitaabar, Justin Chow, Carl Ka To Ma, Barbara J. Ericson, David Weinrop, and Tovi Grossman. 2023. Studying the Effect of AI Code Generators on Supporting Novice Learners in Introductory Programming. *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (2023), 1–23.
- [12] Kenneth R. Koedinger, Albert T. Corbett, and Charles Perfetti. 2012. The Knowledge-Learning-Instruction Framework: Bridging the Science-Practice Chasm to Enhance Robust Student Learning. *Cognitive Science* 36, 5 (2012), 757–798.
- [13] B. Ma, H. Li, G. Li, L. Chen, C. Tang, Y. Xie, C. Gu, A. Shimada, and S. Konomi. 2025. Scaffolding Metacognition in Programming Education. *arXiv preprint arXiv:2511.04144* (2025).
- [14] Hussein Mozannar, Gagan Bansal, Adam Fournier, and Eric Horvitz. 2024. Reading Between the Lines: Modeling User Behavior and Costs in AI-Assisted Programming. *arXiv preprint arXiv:2210.14306* (2024).
- [15] Raja Parasuraman and Victor Riley. 1997. Humans and Automation: Use, Misuse, Disuse, Abuse. *Human Factors* 39, 2 (1997), 230–253.
- [16] Sida Peng, Eirini Kalliamvakou, Peter Cihon, and Mert Demirer. 2023. The Impact of AI on Developer Productivity: Evidence from GitHub Copilot. *arXiv preprint arXiv:2302.06590* (2023).
- [17] James Prather, Brett A. Becker, Michelle Craig, Paul Denny, Dastyni Loksa, and Lauren Margulieux. 2024. The Widening Gap: The Effects of AI-Assisted Code Generation on Novice and Expert Developers. *Proceedings of the 55th ACM Technical Symposium on Computer Science Education* (2024), 142–148.
- [18] Zheyuan Shen, Nikolas Zolas, Samuel Assefa, Miranda Bogen, and Noam Slonim. 2026. How AI Impacts Skill Formation. *arXiv preprint arXiv:2601.20245* (2026).
- [19] M. I. H. Shihab, C. Hundhausen, A. Tariq, S. Haque, Y. Qiao, and B. Mulanda. 2025. The Effects of GitHub Copilot on Computing Students. In *ICER*.
- [20] TUM Applied Education Technologies. 2025. Less Stress, Better Scores, Same Learning. *Computers and Education: Artificial Intelligence* (2025).
- [21] Priyan Vaithilingam, Tianyi Zhang, and Elena L. Glassman. 2022. Expectation vs. Experience: Evaluating the Usability of Code Generation Tools Powered by Large Language Models. *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (2022), 1–7.
- [22] M. Vukovic, R. Pan, T.K. Ho, R. Krishna, R. Pavuluri, and M. Merler. 2026. Usage, Effects and Requirements for AI Coding Assistants in the Enterprise. In *LLM4Code Workshop, ICSE*.
- [23] Celina Weber, Linwei Fang, David Brioneske, and Gunter Saake. 2025. The Impact of AI Tools on Software Engineering Tasks. *arXiv preprint arXiv:2507.09089* (2025).